



# Computer Networks

## Chapter 2 - Getting Connected

Prof. Marino Miculan





## Problems

- In Chapter 1 we saw networks consists of links interconnecting nodes. How to connect two nodes together?
- We also introduced the concept of “cloud” abstractions to represent a network without revealing its internal complexities. How to connect a host to a cloud?



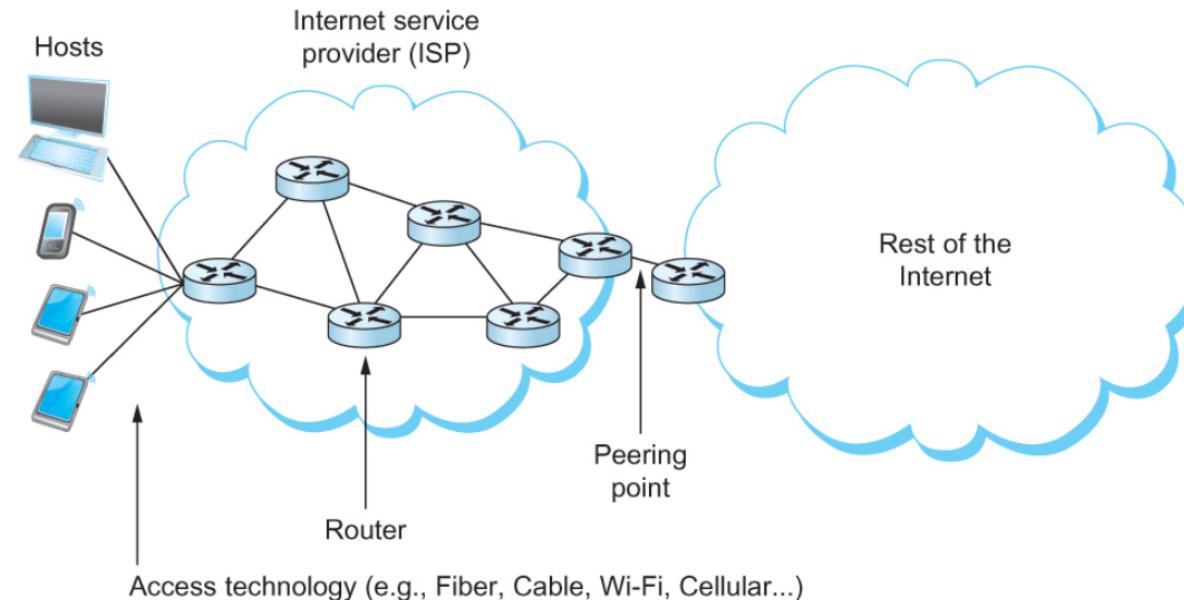
# Chapter Outline

- Perspectives on Connecting nodes
- Encoding
- Framing
- Error Detection
- Reliable Transmission
- Ethernet and Multiple Access Networks
- Wireless Networks

## Chapter Goal

- Exploring different communication medium over which we can send data
- Understanding the issue of encoding bits onto transmission medium so that they can be understood by the receiving end
- Discussing the matter of delineating the sequence of bits transmitted over the link into complete messages that can be delivered to the end node
- Discussing different technique to detect transmission errors and take the appropriate action
- Discussing the issue of making the links reliable in spite of transmission problems
- Introducing Media Access Control Problem
- Introducing Carrier Sense Multiple Access (CSMA) networks
- Introducing Wireless Networks with different available technologies and protocol

# Perspectives on Connecting



An end-user's view of the Internet

# Links

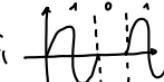
- All practical links rely on some sort of electromagnetic radiation propagating through a medium or, in some cases, through free space
- One way to characterize links, then, is by the medium they use
  - Typically **copper wire** in some form (as in Digital Subscriber Line (DSL) and coaxial cable), → RAME
  - **Optical fiber** (as in both commercial fiber-to-the-home services and many long-distance links in the Internet's backbone), or
  - **Air/free space** (for wireless links) → ACCESSO SATELLITARE

MEZZI per i collegamenti

Le informazioni possono essere trasmesse via cavo variando alcune proprietà fisiche, es. TENSIONE o CORRENTE

## Links

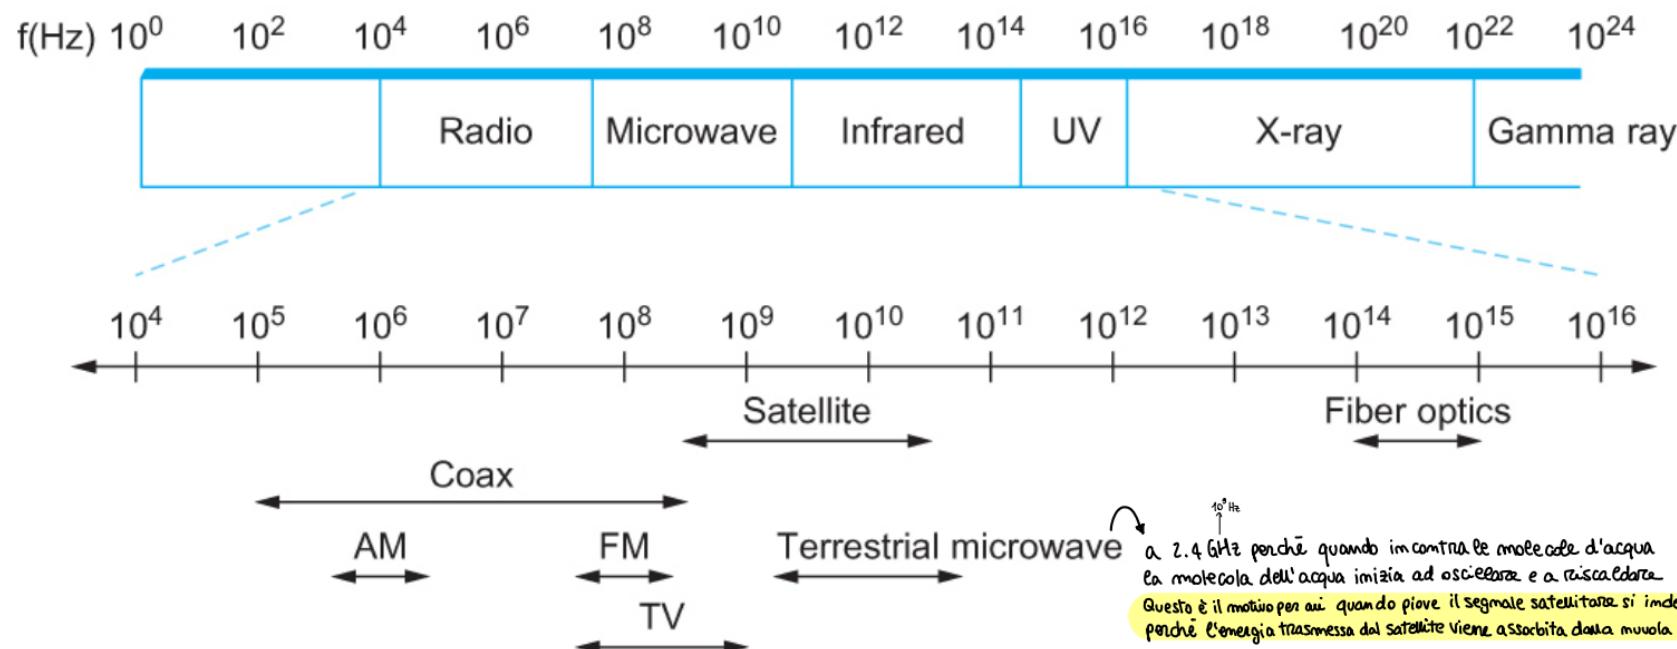
Un mezzo ha un range di frequenze che lo possono attraversare, es. la luce non attraversa i muri

- Another important link characteristic is the **frequency**
  - Measured in hertz, with which the electromagnetic waves oscillate
- Distance between the adjacent pair of maxima or minima of a wave measured in meters is called **wavelength**
- Speed of light divided by frequency gives the wavelength.
  - Ex. (phone communication): **Frequency on a copper cable range from 300Hz to 3300Hz**; Wavelength for 300Hz wave through copper is:  
$$\text{speed of light on copper} / \text{frequency} = \text{about } 2 \times 10^8 / 300 = 667 \times 10^3 \text{ meters.}$$
  - (in fact, the speed depends of the kind of cable)
- **Placing binary data on a signal is called **encoding**.** → codifica dei bit, trasmettersi con delle onde 
  - MODULARE, cambiare le segnale
- **Modulation** involves **modifying the signals in terms of their frequency, amplitude, and phase** (B unchanged!)
- Ex.: shifting a 0-3000 Hz voice signal into phone link frequency range 300-3300 Hz.

# Links

più aumenta la frequenza più è debole verso gli ostacoli

- Electromagnetic spectrum



# Links

| Service                     | Bandwidth (typical) |
|-----------------------------|---------------------|
| Dial-up                     | 28–56 kbps          |
| ISDN                        | 64–128 kbps         |
| DSL                         | 128 kbps–100 Mbps   |
| CATV (cable TV) → COASSIALE | 1–40 Mbps           |
| FTTH (fibre to the home)    | 50 Mbps–1 Gbps      |

throughput

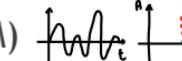
Common services available to connect your home

# Encoding - Signals and Symbols

→ come codificare i bit sul canale

- But how much data can we encode and transmit on a channel?
- There are two aspects to consider:
  - How many bits we can stuff in a **symbol**
  - How many **symbols** we can transmit in a **second**
- A **symbol** is any state in which the link can be set by transmitter, and observed by receiver
- Symbols are what are actually transmitted on the channel. Examples:

- Voltage levels (like in circuits) → es il filo può essere a 0V o 5V

- Amplitude of a carrier frequency (AM)  
 (la frequenza è la stessa, varia l'ampiezza)

- Frequency shift (FM)  
 (la ampiezza è la stessa, varia frequenza)

- Phase shifts (PSK). Very used for digital encoding

- On/off light

- Letters on the paper. Each letter + space is the alphabet

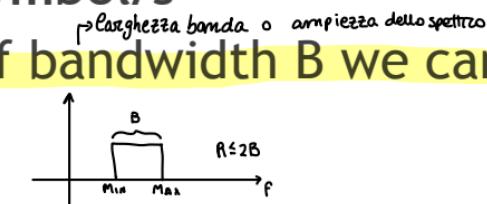
- Language phonemes in the voice

- Smoke puffs, etc.

⇒ La trasmissione di segnali analogici è ottenuta variando la TENSIONE nel tempo, per rappresentare un flusso di informazioni

## Encoding - Symbols and Nyquist-Shannon theorem

- A symbol is any state in which the link can be set by transmitter, and observed by receiver
- The faster we can change the state of the line, the faster we can send symbols, the more data we can transmit over the channel
  - ↳ #max simboli in un secondo
- The number of symbols per second is called **baud rate**, usually denoted with **R**, and measured in **baud = symbol/s**
- **Nyquist-Shannon theorem**: on a channel of bandwidth **B** we can transmit symbols at rate no more than **2B**:  
$$R = 2B$$
  - ↳ max numero di simboli al secondo
- (actually this is a corollary of the real N-S theorem, which says that with a sampling frequency **R** we can correctly reconstruct a signal of frequency **B** up to **R/2**)



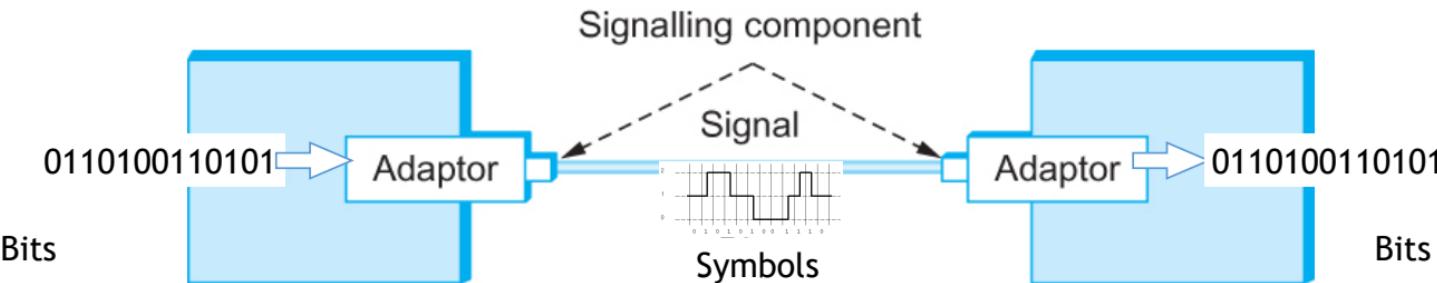
Una linea a n baud trasmette n simboli al secondo e un BAUD può contenere (essere riconosciuto) più bit es. un simbolo è 6 bit

$$\Rightarrow \text{bitrate} = 14,4 \text{ Kbit/s}$$

$$\text{baud rate} = \frac{14,4 \cdot 10^3}{6} = 2,4 \cdot 10^3 \text{ baud/s}$$

## Encoding - Signals and Symbols

- If we have an alphabet of  $M$  symbols, the number of bits that each symbol can represent is  $\log_2(M)$
- For instance, if we encode bits with an electric tension over a wire, within a given voltage range, we can divide this range into
  - 2 levels => 2 symbols = 1 bit per symbol
  - 4 levels => 4 symbols = 2 bits per symbol
  - 128 levels => 128 symbols = 7 bits per symbol
- To transmit a tuple of bits, the encoder signals the symbol (i.e., sets the tension) corresponding to the given tuple. The decoder does the converse.



## Encoding - Symbols

→ il computer lavora in base 2 e per inviarlo in segnali "divide" la tensione in vari intervalli (in base a quale base mi conviene, es. 8 o 16) e una volta ricevuto viene riconvertito in binario.  
⇒ lo faccio per risparmiare simboli

- Example: a range  $-V$  –  $+V$  divided in 8 levels → associa ad ogni valore di tensione 3 bit

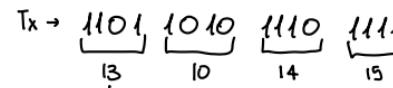
- Symbols are the 8 possible voltage levels; each carries 3 bits

- On an ideal channel, we could divide the signal range as much as we want, and stuff as many bits as we want on one symbol

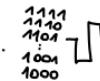
→ si ma non è detto che il bit che manda sia quello che arriva, a causa di DISTURBI E INTERFERENZE (infatti più "suddiviso" la tensione più

rischio errori nella trasmissione perché vado nella tacchetta adiacente)

- Infinite capacity for everyone!

$T_x \rightarrow$  

→ non vengono inviati sequenzialmente, bensì istantaneamente con i LIVELLI es.



Converti in  $b_{16}$ :

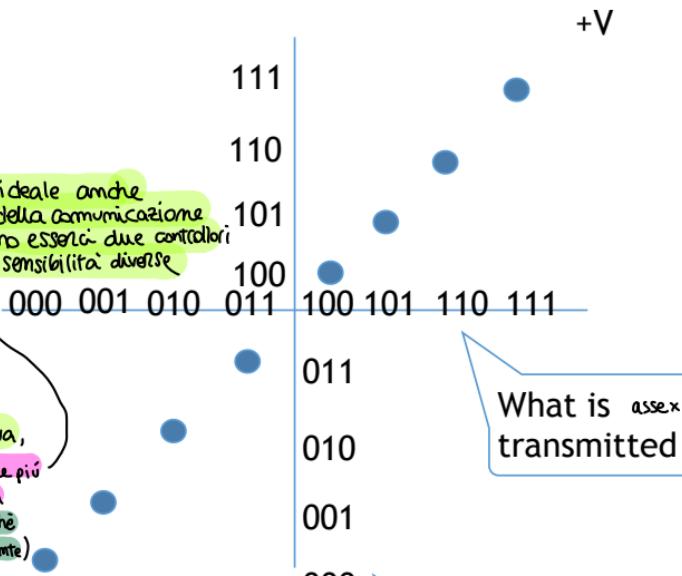
▷

A

E

F

e uso solo 4 simboli al posto che 16, poi in RX converti da  $b_{16}$  a  $b_2$



## Encoding - Noise on the channel

- But real channels have **NOISE**, which is a *random* signal added to the good one
  - Thermal noise on cables, distortions, interferences, technological limits on the precision of encoders and the discriminating power of decoders, etc.
- If the noise is too powerful, a symbol can be changed in the channel, and then misinterpreted by the decoder
  - eg. The tension can increase enough to fall in the next one
- The **important factor** is the **signal-to-noise ratio, SNR**: the ratio between the **average signal power S** and the **average noise power N**

$$\text{SNR} = S/N \quad [\text{pure number}]$$

## Encoding - Noise on the channel

- Often SNR is expressed in **decibel** (dB):

$$\text{SNR} = 10 \cdot \log_{10}(\text{S}/\text{N}) \quad [\text{dB}]$$

- We sometimes say that this is the “noise level”.
- An increment of 3 dB corresponds to about doubling the S/N
- An increment of 10 dB corresponds to increasing the S/N ten times
- For instance
  - SNR=0 dB means  $\text{S}/\text{N} = 1$ , so  $\text{S}=\text{N}$
  - SNR=10 dB means  $\text{S}/\text{N} = 10$
  - SNR=20 dB means  $\text{S}/\text{N} = 100$
  - SNR=30 dB means  $\text{S}/\text{N} = 1000$
  - SNR= -20 dB means  $\text{S}/\text{N} = 0.01$

# Encoding - Link Capacity and Shannon-Hartley Theorem

- The famous **Shannon-Hartley theorem** gives the **upper bound** to the **capacity** of a link in terms of **bits per second** (bps) as a function of **SNR** of the link and the available bandwidth **B** (in Hz).

$$C = B \log_2(1+S/N) = B \log_2(1+SNR) \quad [\text{bps}]$$

↳ il 2 per passare da Hz a bit       $\frac{S}{N}$

Se il segnale è debole rispetto al rumore allora  $SNR \rightarrow 0$  e quindi  $B \cdot \log_2(1+0) \rightarrow 0$ , non riusciamo a trasmettere indipendentemente dalla banda

- Ex: phone link: le frequenze vanno da 300 a 3300 Hz  $\Rightarrow$  perciò la banda (RANGE) è 3000 Hz.
  - $B = 3300 - 300 = 3000$  Hz, S is the (average) signal power, N the average noise power.  
perchè ricordiamo che  $SNR = 10 \cdot \log_{10}(S/N)$  [dB]
  - If there is 30 dB of noise then  $\overset{\text{SNR}}{S/N} = 1000$ .  
anche se la banda è piccola ma c'è poco rumore, io posso trasmettere tanta informazione
  - Now  $C = 3000 \times \log_2(1001) = 30$  kbps.
- To get a larger capacity, we need to **improve the bandwidth**, or the **quality of the line (SNR)**, or both

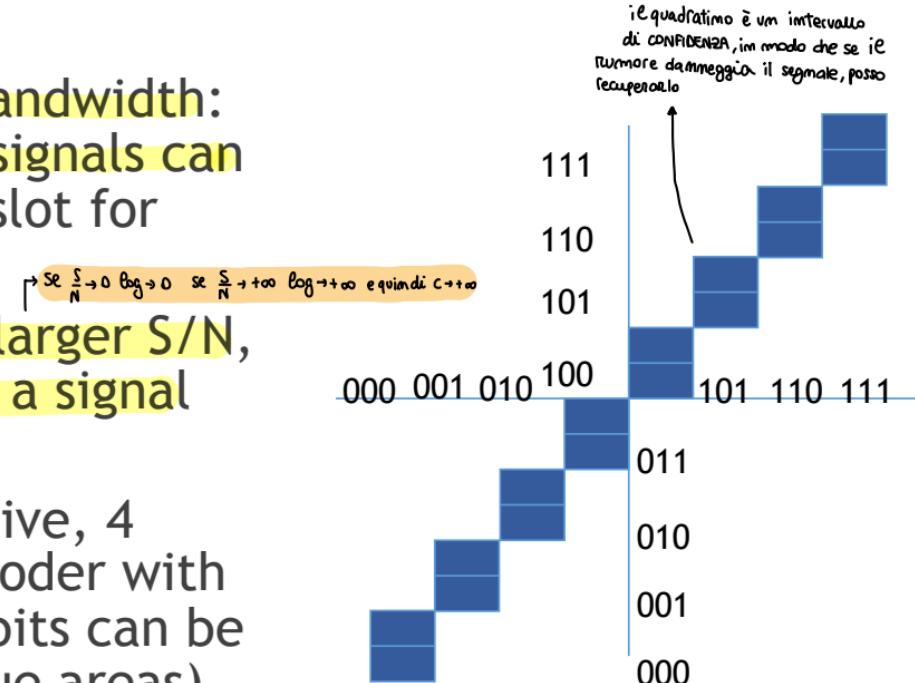
→ se c'è troppo rumore devo diminuire i simboli, se ce n'è poco posso aumentarli

↳ nel calcolo è  $B \cdot \log_2(1 + \frac{S}{N})$

BANDA      QUALITÀ DELLA RETE

# Encoding: Link Capacity and Shannon-Hartley Theorem

- Capacity is proportional to the bandwidth: the larger  $B$ , the faster the rate signals can be commuted (i.e. smaller time slot for sending a symbol)
- Dependency on  $\log_2(1+S/N)$ : the larger  $S/N$ , the more bits can be encoded on a signal level.
- Ex:  $N=S/7 \Rightarrow$  eight levels (4 positive, 4 negative) can be read by the decoder with no noise interference. Hence, 3 bits can be encoded on each level (noise: blue areas)



- 15 fomemi/s e ci sono 45 fomemi

qual è il throughput  $\rightarrow \text{bit/fomema} = \log_2(45) \rightarrow \text{bit che rappresentano 1 fomema} = 5,49 \rightarrow \text{per codificare 1 fomema uso 5,5 bit}$

$$\text{ie bit-rate } C = 5,49 \cdot 15 = 82,3 \text{ bit/s}$$

$$\left[ \frac{\text{bit}}{\text{fomema}} \right] \left[ \frac{\text{fomema}}{\text{s}} \right]$$

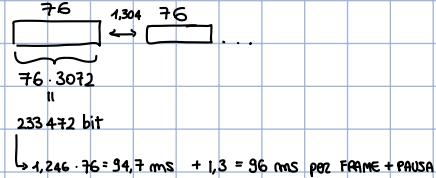
Se SNR = 10 dB, qual è l'ampiezza di banda?

$$B = ? \quad \text{SNR} = 10 \text{ dB} = 10 \log_{10} \left( \frac{S}{N} \right) \Rightarrow 10 = 10 \log_{10} \left( \frac{S}{N} \right) \Rightarrow 1 = \log_{10} \left( \frac{S}{N} \right) \Rightarrow \frac{S}{N} = 10$$

$$\Rightarrow 82,3 = B \cdot \log_2 (1+10) \Rightarrow B = \frac{82,3}{\log_2(11)} = 23,8 \text{ Hz}$$

- ogni simbolo porta 3072 bit e dura 1,246 ms

FPAME ha 76 simboli e i frame sono separati da una pausa di 1,304 ms



$$C = \frac{233.472}{96} = 2,4 \cdot 10^6 \text{ bit/s} = 2,4 \text{ Mb/s}$$

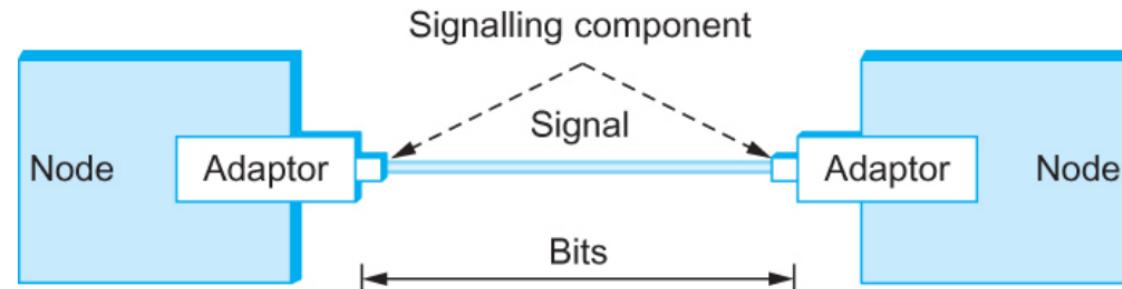
$B = 1536 \text{ KHz}$  (anche se la dab ha banda 160 Hz, me usa un piccolo intervallo)

$$C = B \cdot \log_2 \left( 1 + \frac{S}{N} \right)$$

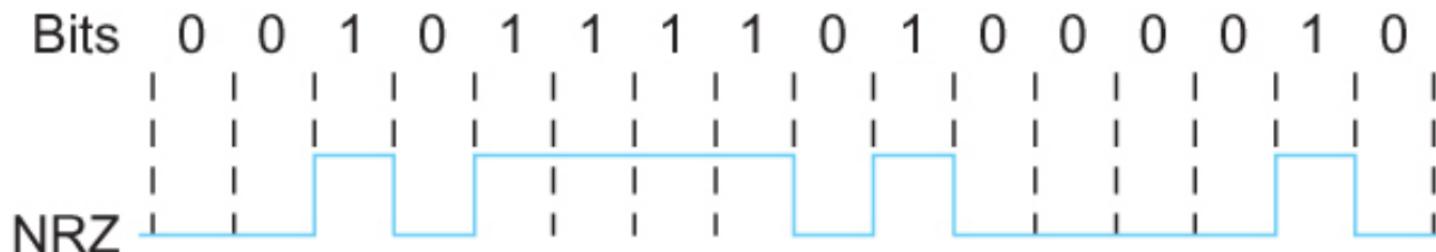
$$1,536 \cdot 10^6 \cdot \log_2 \left( 1 + \frac{S}{N} \right)$$

$$\log_2 \left( 1 + \frac{S}{N} \right) = \frac{2,4}{1,536} \Rightarrow 1 + \frac{S}{N} = 2^{\frac{2,4}{1,536}} \Rightarrow \frac{S}{N} = 2^{\frac{2,4}{1,536}} - 1$$

## Encoding



Signals travel between signalling components; bits flow between adaptors  
↳ alle estremità



NRZ encoding of a bit stream → i dati vengono passati direttamente come tali: in uscita

# Encoding - Problems with NRZ

## • Baseline wander

- The receiver keeps an **average of the signals** it has seen so far
- **Uses the average to distinguish between low and high signal**
- When a signal is significantly low than the average, it is 0, else it is 1
- Too **many consecutive 0's and 1's cause this average to change, making it difficult to detect**

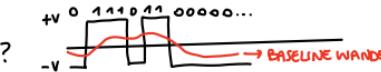
## • Clock recovery

→ difficilissimo tenere sincronizzati due clock

- Frequent transition from high to low or vice versa are necessary to enable **clock recovery** → quando hai tanti "su e giù" le clock si risintonizzano, se hai lunghi periodi statici le clock tende a pendersi
- Both the **sending and decoding process is driven by a clock**
- Every clock cycle, the sender transmits a bit and the receiver recovers a bit
- **The sender and receiver have to be precisely synchronized**

come si fa a capire se è alto o basso?

→ VALORE MEDIO



→

Lo fa per capire più che altro come gestire con i picchi alti per distinguere poi quelli più bassi.

Tempo il valore di MEDIA e se un Segnale è tanto più basso della media allora è 0, se è più alto è 1

→ La media si alza o abbassa troppo e rischia di confondere valori. es. se la media → 1 e arriva 0,88 rischia, essendo sotto la media, di approssimare a 0.

# Encoding - NRZI and Manchester

- **NRZI: Non Return to Zero Inverted**

- Sender makes a transition from the current signal to encode 1 and stay at the current signal to encode 0
- Solves for consecutive 1's, ma non quello degli 0 consecutivi (perché con lo 0 è statico) → - Soffre di BASELINE WANDER (solo per lo 0, ma 1)  
- Soffre di clock-sync

- **Manchester encoding** → non da problemi con lunghe sequenze di 0 e 1 (quindi non perde sincronismo)

- Merging the clock with signal by transmitting Ex-OR of the NRZ encoded data and the clock
- Clock is an internal signal that alternates from low to high, a low/high pair is considered as one clock cycle
- In Manchester encoding
  - 0: low to high transition
  - 1: high to low transition

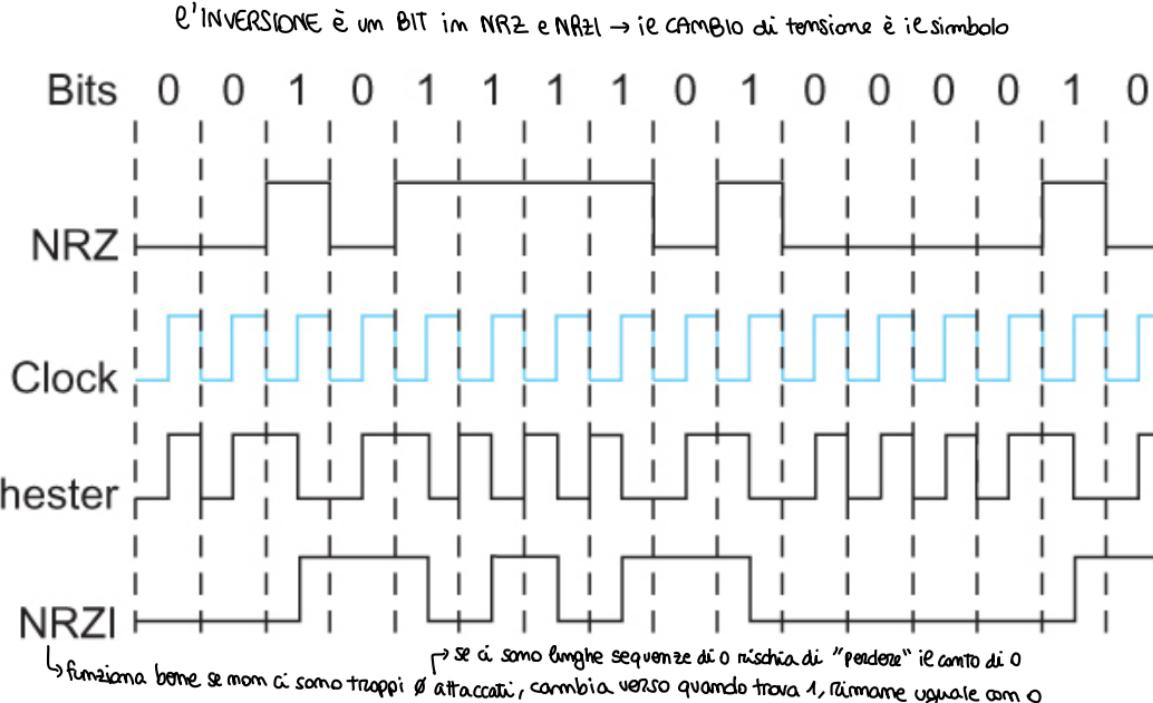
- NON soffre di BASELINE WANDER

- NON soffre di CLOCK-SYNC

- DIMEZZA la BANDA perché in un CICLO DI CLOCK il ricevitore verifica 2 volte la direzione della salita/discesa in 1 istante

# Encoding - different encodings

si prende il clock e lo  
si manda se è 0, immediatamente se è 1  
è autocorrente, il ricevitore si sincronizza  
sul suo clock. I simboli sono i cambi di tensione



Different encoding strategies

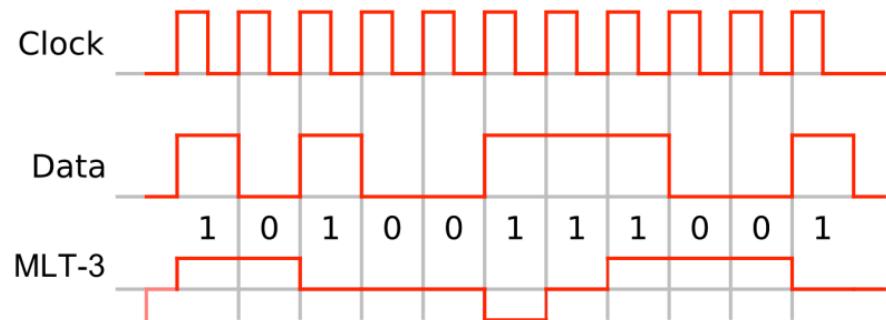
# Encoding

- Problem with Manchester encoding
  - Doubles the rate at which the signal transitions are made on the link
    - Which means the receiver has half of the time to detect each pulse of the signal
  - The rate at which the signal changes state is called the link's **baud rate**
  - If we consider the levels as state, in Manchester we need 2 states to represent one bit (low→high=0, high→low=1) thus **the bit rate is half the baud rate** (2 states = 1 bit)

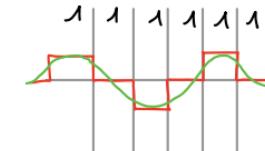
## Encoding - MLT-3

→ Ethernet si basa su questo principio

- MultiLine Transmission, with 3 levels: very similar to NRZI, but it uses 3 levels of signal: -1, 0, +1



Se trasmetto solo 1111...



Se invece è NRZI con tutti 1 ottengo un clock

- Signal changes from one level to the next at the beginning of a “1” bit. No transition at the beginning of “0” bit
- When it reaches one extreme (1 or -1), it changes direction.
- Risk of losing synchronisation for long sequences of “0” → perché non cambia mai verso

## Encoding - 4B/5B encoding → rompe le lunghe sequenze di 0

- Insert extra bits into bit stream so as to break up the long sequence of 0's and 1's
- Every 4-bits of actual data are encoded in a 5- bit code that is transmitted to the receiver → in modo che non puoi mai avere più di 3 zeri consecutivi (guarda la tabella.)
- 5-bit codes are selected in such a way that *each one has no more than one leading 0 (zero) and no more than two trailing 0's.*
- In this way, three consecutive 0's can never appear, even between two different 5-bit codes.
- Efficiency: 80%
- Delay: 5 bits (receiver needs to get 5 bits to decode 4 bits).
- Used often in conjunction with MLT-3 (Ethernet 100 Mbit)

# Encoding - 4B/5B encoding

- 16 codes are left out from encoding
- 7 codes violate the rule *no more than one leading 0 (zero) and no more than two trailing 0's*.

In particular:

- 11111 - when the line is idle
  - quando non sto trasmettendo niente
- 00000 - when the line is dead
  - la linea è scattata
- 6 of these are used for control
  - 00100 - to mean halt
  - 00111 - reset
  - ...

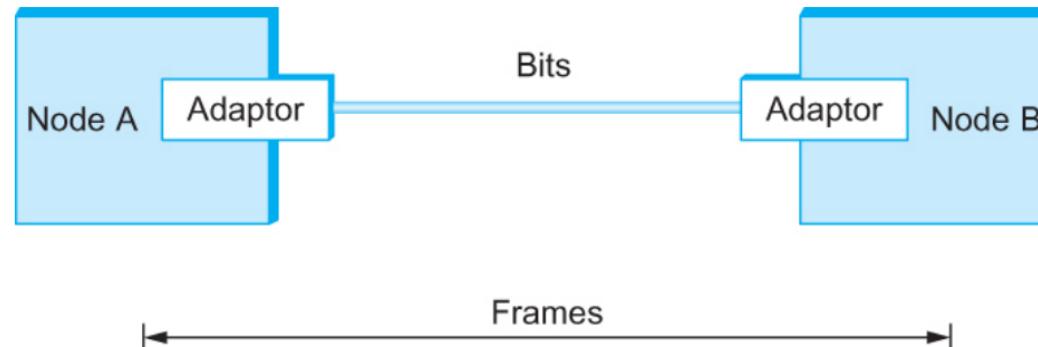
FAST ETHERNET (100 Mb/s)  
Usa 4B15B, in realtà trasmessi 125 Mb/s

| <i>Data Sequence</i> | <i>Encoded Sequence</i> | <i>Control Sequence</i> | <i>Encoded Sequence</i> |
|----------------------|-------------------------|-------------------------|-------------------------|
| 0000                 | 11110                   | Q (Quiet)               | 00000                   |
| 0001                 | 01001                   | I (Idle)                | 11111                   |
| 0010                 | 10100                   | H (Halt)                | 00100                   |
| 0011                 | 10101                   | J (Start delimiter)     | 11000                   |
| 0100                 | 01010                   | K (Start delimiter)     | 10001                   |
| 0101                 | 01011                   | T (End delimiter)       | 01101                   |
| 0110                 | 01110                   | S (Set)                 | 11001                   |
| 0111                 | 01111                   | R (Reset)               | 00111                   |
| 1000                 | 10010                   |                         |                         |
| 1001                 | 10011                   |                         |                         |
| 1010                 | 10110                   |                         |                         |
| 1011                 | 10111                   |                         |                         |
| 1100                 | 11010                   |                         |                         |
| 1101                 | 11011                   |                         |                         |
| 1110                 | 11100                   |                         |                         |
| 1111                 | 11101                   |                         |                         |

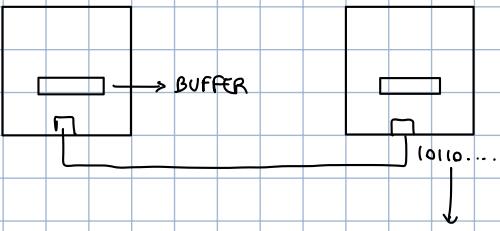
ORA SIAMO nel livello 2      BIT → FRAME + controlli d'errore

**Framing** → d'ora in poi pensiamo di trasmettere bit, escludiamo i dettagli elettrici

- We are focusing on **packet-switched networks**, which means **that blocks of data** (called **frames** at this level), not bit streams, are **exchanged between nodes**.
- It is the network adaptor that enables the nodes to exchange frames.



Bits flow between adaptors, frames between hosts



| FRAME |      |   |
|-------|------|---|
| H     | DATA | T |

- all'interno di questa sequenza
1. deve riconoscere dove inizia e dove finisce il frame ovvero frame header + data + data link trailer
  2. vedere se questi dati sono corretti (se bit sono stati alterati)

⇒ in ricezione raggruppo i bit (che sono un flusso continuo) in byte e poi, ad es. con i caratteri sentinel, inizia a riconoscere il frame

## Framing

- When node A wishes to transmit a frame to node B, it tells its adaptor to transmit a frame from the node's memory. This results in a sequence of bits being sent over the link (which can be point-to-point or shared).
- The adaptor on node B then collects together the sequence of bits arriving on the link and deposits the corresponding frame in B's memory.
- Recognizing exactly what set of bits constitute a frame—that is, determining where the frame begins and ends—is the central challenge faced by the adaptor

# Framing

- **Byte-oriented Protocols**

- To view each frame as a collection of bytes (characters) rather than bits
- **BISYNC** (Binary Synchronous Communication) Protocol
  - Sentinel-based → caratteri sentinella che indicano l'inizio e la fine
  - Developed by IBM (late 1960)
- **DDCMP** (Digital Data Communication Protocol)
  - Byte-count based
  - Used in DECNet
- (but other variants are possible)

## Framing

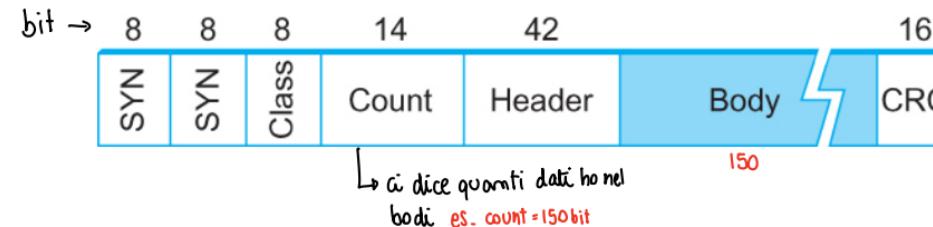
→ caratteri sentinel che indicano l'inizio e la fine

- **BISYNC - sentinel approach** (lavoro fatto in fase di decodifica)
  - Frames transmitted beginning with leftmost field
  - Beginning of a frame is denoted by sending a special SYN (synchronize) character
  - Data portion of the frame is contained between special sentinel character STX (start of text) and ETX (end of text)
    - SOH : Start of Header
    - DLE : Data Link Escape (“escaping”): Body vs. ETX
    - CRC: Cyclic Redundancy Check → trailer



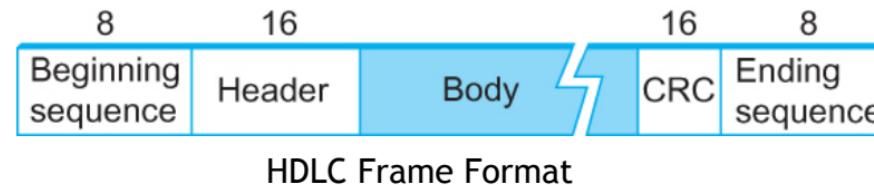
## Framing

- Byte-counting approach: DDCMP
  - *count*: how many bytes are contained in the frame body
  - CRC: Cyclic Redundancy Check as before
  - If count is corrupted: Framing error



## Framing

- Bit-oriented Protocol: frames as sequences of bits
  - Beginning and Ending Sequences
  - HDLC (IBM): High Level Data Link Control
    - Delimiting sequence: 0 1 1 1 1 1 1 0



- And what if delimiting sequence occurs within the frame?

## Framing

- HDLC Protocol, **on the sending side**:
  - **any time five consecutive 1's have been transmitted** from the body of the message (i.e. excluding when the sender is trying to send the distinguished 01111110 sequence), **the sender inserts 0 before transmitting the next bit** ("bit stuffing")
  - Eg: 011111xy becomes 0111110xy

Rompo le sequenze di 1 (nel body) per evitare che vengano scambiati per delimitatori (01111110)

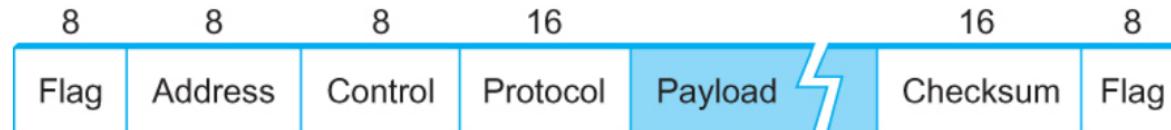
es. 011111 e qui metto uno 0 poi il valore che c'era (es. 1) e ci aggiungo un altro 0 => così non lo scambio per delimitatore

## Framing

- HDLC Protocol, on the **receiving side**:
- **After** having received **5 consecutive 1's**, read **next bit**
  - **If 0** : Stuffed, so discard it and proceed
  - **If 1** : **Either End of the frame marker, or Error has been introduced in the bitstream.** To tell, look at the next bit
    - **If 0** ( 01111110 ) then **End-of-frame marker** => frame ended
    - **If 1** ( 01111111 ) then **Error** => discard the whole frame  
The receiver needs to wait for next 01111110 before it can start receiving again

## Framing

- **PPP (Point-to-Point Protocol) is commonly run over Internet links**



- Derived from HDLC, **uses sentinel approach and bit stuffing**
  - Special start of text character denoted as **Flag**
    - 0 1 1 1 1 1 1 0
  - Address, control : default numbers
  - Protocol for demux : IP / IPX / ...
  - Payload : negotiated by LCP (default: 1500 bytes)
  - Checksum : for error detection (CRC-16 or CRC-32)

→ meglio avere frame grandi per ridurre overhead

meglio avere frame piccoli per ridurre la propria d'ordine

# Error Detection

- Errors are introduced into frames due to *electrical interference* and *thermal noises*
  - The former gives rise to (deterministic) bursts -> sequences of bits destroyed
  - The latter to rare, supposedly independent one-bit spikes
- Detecting errors is a fundamental function of the datalink layer
- Approaches when the recipient detects an error:
  - **Drop silently** the frame - as it was never sent → *il più semplice*
    - Upper layers may recognise the packet loss, and take suitable countermeasures (if they need it)
  - **Notify the sender** that the message was corrupted, so the sender can send again.
    - If the error is rare, then the retransmitted message will (likely) be error-free
  - **Forward Error Correction:** Using some **error correction algorithm**, the receiver reconstructs the message even in presence of errors (up to some extent)
    - ↳ il sender oltre a mandare informaz. per capire se è errato può anche mandare codici di connessione → **OVERHEAD**
    - ↳ spesso quando la trasmissione è unidirezionale (es. antenna broadcast sky) dove non posso chiedere di ricontraddomi i pacchetti

## Bit errors: simple examples

$$\gamma(A \cup B)$$

$\Downarrow$

$$\gamma A \wedge \gamma B$$

- Joint occurrence of two rare events A and B:

$$\begin{aligned} P[A \cup B] &= 1 - P[\downarrow(A \cup B)] = 1 - P[(\sim A) \cap (\sim B)] \\ &= 1 - (1-p_A)(1-p_B) = p_A + p_B - p_A p_B \xrightarrow[\text{tende a 0}]{\text{sono eventi rari,}} \approx p_A + p_B \\ &\xrightarrow{\text{ALMENO UNO}} \end{aligned}$$

- At least one error in a 10 kb packet,  $p_E \approx 10^{-7}$ , bit independence approximation:

$$\begin{aligned} P[\{\text{one error in 10 kb}\}] &= 1 - P[\{\text{no error in 10 kb}\}] = 1 - (p_{\text{bit corretti}})^n \\ &\xrightarrow{\text{almeno un errore}} = 1 - (1-10^{-7})^{10000} = 0.0009995 \approx 10^4 \cdot 10^{-7} = 10^4 \cdot p_E \end{aligned}$$

- Two errors in the same packet, same error, same approximation:

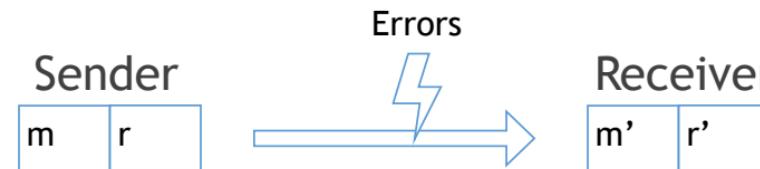
$$\begin{aligned} P[\{\text{two errors in 10 kb}\}] &\approx (1-p_E)^{10000-2} \cdot (p_E)^2 \cdot C(10^4, 2) = \\ &(1-10^{-7})^{9998} \cdot 10^{-14} \cdot 10^4 \cdot (10^4-1)/2 \approx 10^{-14} \cdot 5 \cdot 10^7 = 5 \cdot p_E \end{aligned}$$

# Error Detection

- Basic Idea of Error Detection:
  - To add redundant information to a frame that can be used to determine if errors have been introduced
- Extreme Case: transmitting two complete copies of data
  - Identical → No error
  - Differ → Error
- Poor Scheme?
  - n bit message, n bit redundant information
  - Error can go undetected → se un errore è presente nella stessa posizione in entrambi non verrà rilevato
- In general, we can provide strong error detection technique
  - n bits message, k redundant bits
  - Usually k is constant, and  $\ll n$
  - (In Ethernet, a frame carrying up to 12,000 bits of data requires only 32-bit CRC)

## Error Detection

- Extra  $k$  bits are redundant: They add no new information to the message
- Derived from the original message using some algorithm
- Both the sender and receiver know the algorithm
- Result of the transmission:



- Receiver computes  $r''$  using  $m'$  (and the same algorithm)
  - if  $r' = r'' \rightarrow$  no error
  - if  $r'$  differs from  $r'' \rightarrow$  error

# Error Detection

- **CRC (Cyclic Redundancy Check)**
  - Common technique for detecting transmission error
  - Used in DDCMP, HDLC, CSMA/CD, Token Ring, etc.
- Other approaches
  - **Parity**
  - **Two Dimensional Parity (BISYNC)**
  - **Checksum (IP, UDP, TCP)**

## One-dimensional parity

→ aggiungere 1 bit di parità

- “simple” (one-dimensional) parity involves adding one extra bit to a n-bit code to balance the number of 1s in the byte.
- For example, on 7-bit codes
  - Odd parity sets the eighth bit to 1 if needed to give an odd number of 1s in the byte
  - Even parity sets the eighth bit to 1 if needed to give an even number of 1s in the byte
- Each byte is checked separately from the others
- An even number of errors in a byte goes undetected

→ non rileva modifiche pari di bit

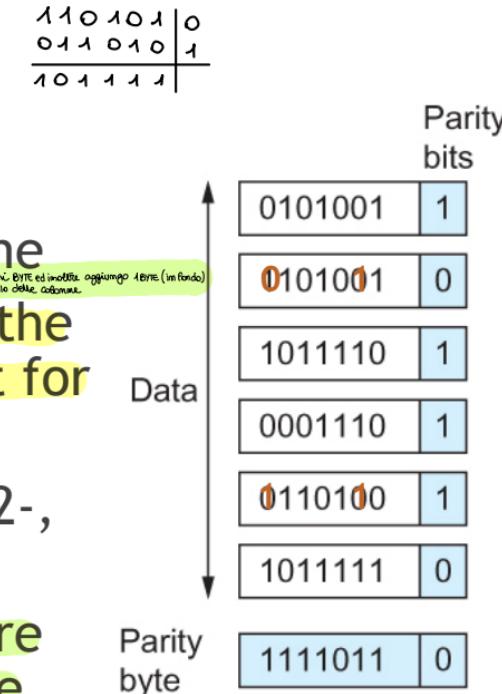
↓  
il frame viene letto  
a byte e quindi ad ogni  
byte viene aggiunto un bit  
di parità

| Parity bits |   |
|-------------|---|
| 0101001     | 1 |
| 1101001     | 0 |
| 1011110     | 1 |
| 0001110     | 1 |
| 0110100     | 1 |
| 1011111     | 0 |

## Two-dimensional parity

anche i bit di parità possono essere modificati

- Two-dimensional parity does a similar calculation for each bit position across each of the bytes contained in the frame
  - This results in an extra *parity byte* for the entire frame, in addition to a parity bit for each byte → controllo i bit di parità sia in orizzontale che verticale
  - Two-dimensional parity catches all 1-, 2-, and 3-bit errors and most 4-bit errors
  - 4-bit errors are not detected if there are “aligned” in same rows AND in the same columns → se si cambiamo 4 bit a quadrato e allora non li rileva



Four bit errors which are not detected by parity bits, nor by two-dimensional parity

## Internet Checksum Algorithm

↳ viene usato a livelli 3-4 ma non a livello 2

- Consider the data as a sequence of 16-bit integers
- Add them together using 16-bit ones-complement arithmetic (explained next slide) and then take the ones complement of the result.
  - The resulting 16-bit number is called the checksum
- Transmit the checksum together with the message
- The receiver performs the same calculation on the received data and compares the result with the received checksum → **dal messaggio si calcola il checksum e se coincide con il checksum che ha ricevuto allora NO ERRORI**
- If any transmitted data, including the checksum itself, is corrupted, then the results will not match, so the receiver knows that an error occurred
- Not used at the link level, but used at network and transport level

## Internet Checksum Algorithm

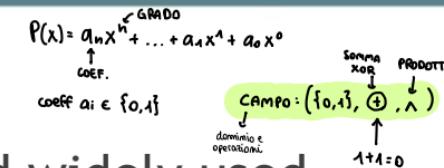
- In ones-complement arithmetic, a negative integer  $-x$  is represented as the complement of  $x$ ; E.g.:  $2 \Leftrightarrow 0010$  iff  $-2 \Leftrightarrow 1101$
- When adding numbers in ones complement arithmetic, a carryout from the most significant bit needs to be added to the result.
  - This happens because zero has two logical representations:  $00..0$  and  $11..1$ : we need to skip the latter when the sum wraps around the represented domain
    - ↗ devo solo riportare tra positivi e negativi
- Example: addition of  $-5$  and  $-2$  in ones complement arithmetic on 4-bit integers
  - $+5$  is  $0101$ , so  $-5$  is  $1010$ ;  $+2$  is  $0010$ , so  $-2$  is  $1101$ 
    - ↗ RIPORTO
  - If we add  $1010$  and  $1101$  ignoring the carry, we get  $0111$
  - Since there is a carry from the most significant bit causes us to increment the result, giving  $1000$ , which is the ones complement representation of  $-7$ , as we would expect
- (The book provides a weird example:  $-8$  cannot be represented in ones complement arithmetic!)

# Internet Checksum Algorithm

```
u_short cksum(u_short *buf, int count) {
    register u_long sum = 0;    \\\ 16-bit checksum
    while (count--) {          \\\ count 16-bit units
        sum += *buf++;        \\\ 16-bit buffer units
        if (sum & 0xFFFF0000) {  \\\ if carry
            sum &= 0xFFFF;      \\\ erase carry
            sum++;              \\\ increment
        }
    }
    return ~(sum & 0xFFFF);  \\\ erase last carry
}                                \\\ and complement
```

## Cyclic Redundancy Check (CRC)

↪ usato ovunque



- Introduced by Wesley Peterson (1961), and widely used
- Reduce the number of extra bits and maximize protection
- Based on Modulo 2 polynomial arithmetics
- Given a bit string we can associate a polynomial  $P(x)$  for it
- E.g. 110001 corresponds to the polynomial of degree 5
$$1 \cdot x^5 + 1 \cdot x^4 + 0 \cdot x^3 + 0 \cdot x^2 + 0 \cdot x^1 + 1 \cdot x^0 = x^5 + x^4 + 1 \rightarrow \text{considero la stringa come somma}$$

$a_n x^n + a_{n-1} x^{n-1} + \dots$   
 $\Downarrow$   
 $a_n a_{n-1} \dots$
- Thus a bit string is just another representation for a polynomial
- A  $k$ -bit frame has a maximum degree of  $k-1$
- Let  $M(x)$  be a message polynomial and  $C(x)$  be a given generator polynomial.

# Cyclic Redundancy Check (CRC)

$$\text{XOR} \rightarrow \begin{array}{r} \text{-} \mid 0 \mid \\ \mid 0 \mid \\ \hline 1 \mid 1 \mid 0 \end{array} \quad \begin{array}{r} \text{+} \mid 0 \mid \\ \mid 0 \mid \\ \hline 1 \mid 0 \mid 0 \end{array}$$

## Polynomial Arithmetic Modulo 2

- Polynomials can be added and multiplied (easy)
- Any polynomial  $B(x)$  can be divided by a divisor polynomial  $C(x)$  if  $B(x)$  is of higher degree than  $C(x)$  (note: true for any Modulo b polynomial arithmetic)
- Any polynomial  $B(x)$  can be divided once by a divisor polynomial  $C(x)$  if  $B(x)$  is of the same degree as  $C(x)$ :  
if  $\deg(B) = \deg(C)$  then  $B(x)/C(x) = 1 + \text{remainder}$
- In this case, the remainder is obtained by subtracting  $C(x)$  from  $B(x)$ :  
 $B(x)/C(x) = 1$  with remainder  $R(x)$  implies  
 $B(x) = 1 \cdot C(x) + R(x)$  and hence  $R(x) = B(x) - C(x)$ .
- To subtract  $C(x)$  from  $B(x)$ , we simply perform the exclusive-OR (XOR) on each pair of matching coefficients (no carry, same as summing):  
 $B(x) + C(x) = B(x) - C(x) = B(x) \text{ XOR } C(x)$

SOMMA:  $P_1(x) + P_2(x) = \sum_{i=0}^n (a_i \oplus b_i) x^i$

$$P_1(x) + P_4(x) = \sum_{i=0}^n (a_i \oplus a_i) x^i = 0 \text{ perché } 1+1=0$$

PRODOTTO:

es.  $P(x) = x^3 + x^2 + 1$   
 $(1101)$   
 $C(x) = x^2 + 1$   
 $(101)$

$$P(x) \cdot C(x) = x^5 + x^3 + x^4 + x^2 + x^2 + 1 = x^5 + x^4 + x^3 + 1$$
  
 $(111001)$

|     |   |     |   |     |
|-----|---|-----|---|-----|
| XOR | - | 0 1 | ⊕ | 0 1 |
|     | 0 | 0 1 | 0 | 0 1 |
|     | 1 | 1 0 | 1 | 1 0 |

SOMMA:  $P_1(x) + P_2(x) = \sum_{i=0}^n (a_i \oplus b_i) x^i$

$$P_1(x) + P_2(x) = \sum_{i=0}^n (a_i \oplus a_i) x^i = 0 \text{ perché } 1+1=0$$

PRODOTTO:

es.  $P(x) = x^3 + x^2 + 1$   
 $(1101)$

$C(x) = x^2 + 1$   $P(x) \cdot C(x) = x^5 + x^3 + x^4 + x^2 + x^2 + 1 = x^5 + x^4 + x^3 + 1$   
 $(101)$   $(111001)$

se  $B(x)$  ha lo stesso grado di  $C(x)$

DIVISIONE:  $B(x) = Q(x) \cdot C(x) + R(x)$  grado  $(R) < \text{grado}(C)$  grado  $(Q) = \text{grado}(B) - \text{grado}(C)$

B 110101 → faccio XOR: R = 011100  $\Rightarrow$  che sarebbe  $\frac{(110101)}{x^5 + x^4 + x^2 + 1} / \frac{(101001)}{x^5 + x^3 + 1}$

C 101001

io mando  $P(x)$  e ricevo  $P'(x)$   
 alcuni bit scappati

$$P'(x) = P(x) + E(x)$$

es. TX: 100101...  
 RX: 110100...  
 E(x) = 010001...

$$\text{resto } (P(x), C(x)) = \text{resto } (E(x), C(x))$$

rilevo gli errori facendo la divisione tra  $E(x)$  e  $C(x) \Rightarrow \frac{E(x)}{C(x)}$

se il resto è 0 è molto probabile che non ci siamo errore  
 resto è 1 allora c'è errore → scarto

NON RILEVO quando  $E(x)$  è multiplo di  $C(x)$

## Cyclic Redundancy Check (CRC)

- Let  $P(x)/C(x)$  leave a remainder  $R(x)=0$ .
- The receiver's viewpoint: error is a sequence of bits, i.e. a polynomial  $E(x)$

- When  $P(x)$  is sent and  $P'(x)$  is received we have

$$P'(x) = P(x) + E(x)$$

resto

- The receiver compute  $R'(x)$ , the remainder of  $P'(x)/C(x)$

- $R'(x) = \text{remainder } [P'(x)/C(x)] = \text{remainder } [E(x)/C(x)]$

→ perché  $P = P+E$  e  $P/C$  ha resto  $\emptyset$ , perciò  $P'/C$  ha resto uguale a  $E/C$

- If it nonzero, then an error (non divisible by  $C(x)$ ) has occurred!

- The only thing the sender and the receiver should know is  $C(x)$ .

## Cyclic Redundancy Check (CRC)

→ ie controllo avviene a posteriori

- Let  $M(x)$  come from a frame with  $m$  bits and
- Let the generator polynomial  $C(x)$  comes from a frame having less than  $m$  bits: say,  $\deg(C) = r$ .

$$P(x) = \overset{\text{messaggio}}{M(x)} \cdot \overset{\text{generatore}}{C(x)}$$

- The **logical algorithm**:

send  $P(x) = M(x)C(x)$  (which is divisible by  $C(x)$ , obviously)

Receivers gets  $P'(x) = P(x) + E(x)$

Then, no errors are detected by the receiver iff

$R'(x) = \text{reminder}[P'(x)/C(x)] = \text{reminder}[E(x)/C(x)] = 0$

- How to perform multiplication  $M(x)C(x)$ ?

# Cyclic Redundancy Check (CRC)

The algorithm *in practice*:

- Append  $r$  zeros to the low-order end of the frame, so it now contains  $m+r$  bits this gives the polynomial  $x^r M(x)$ .  
sto shiftando a sx di  $r$  posizioni moltiplico tutto per  $x^r$
- Divide  $x^r M(x)$  by  $C(x)$  using modulo 2 division (see algorithmic procedure next slide)
- Subtract the remainder  $R(x)$  (which is always  $r$  or fewer bits) from the string corresponding to  $x^r M(x)$  using modulo 2 subtraction (immediate:  $R(x) \text{ XOR } 0 = R(x)$ )
- The result  $P(x) = x^r M(x) - R(x) = x^r M(x) + R(x)$  is the checksummed frame to be transmitted. Notice that  $\text{remainder}[P(x)/C(x)] = 0$

SENDER

- The receiver receives  $P'(x) = P(x) + E(x)$  and computes  $R'(x) = \text{remainder}[P'(x)/C(x)] = \text{remainder}[E(x)/C(x)]$ 
  - if  $R'(x) = 0$  then it discards  $r$  LSB's, i.e. it takes  $M(x) = x^r P'(x)$ ; otherwise there is an error

RECEIVER

# Cyclic Redundancy Check (CRC)

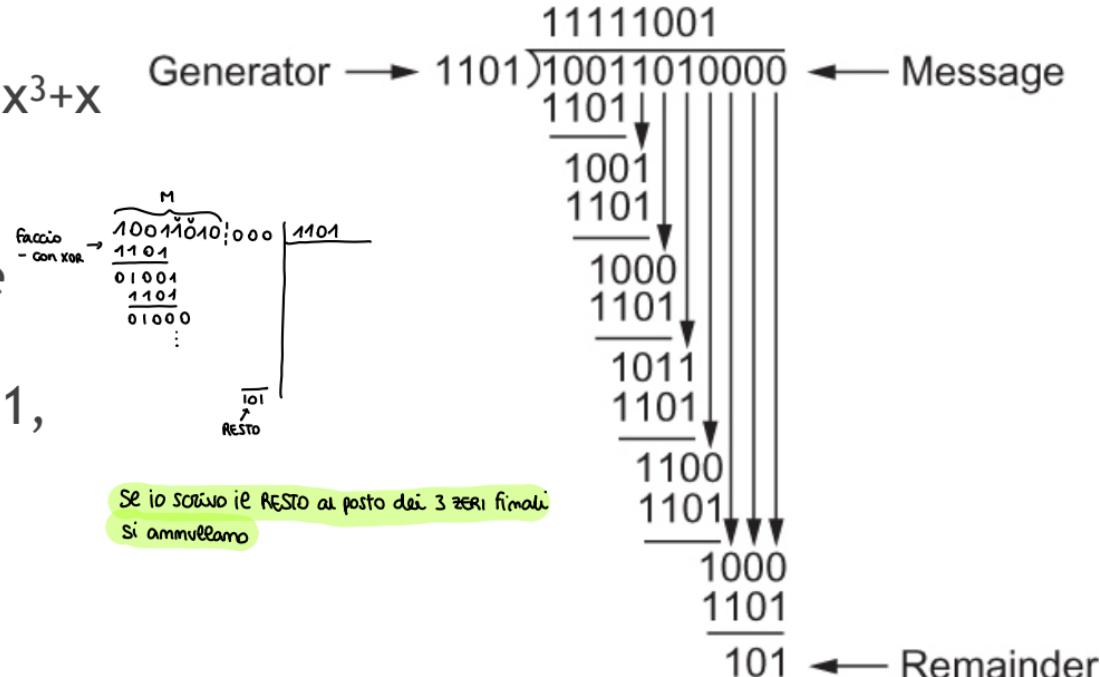
- Example:

$$M = 10011010, \quad M(x) = x^7 + x^4 + x^3 + x$$
$$C = 1101, \quad C(x) = x^3 + x^2 + 1$$

- $\deg(C(x)) = 3$ , so we add 3 zeros at the end of  $M$  (the CRC will have 3 bits)

- The remainder is  $R(x) = x^2 + 1$ , that is  $R = 101$

- $P(x) = x^3M(x) + R(x)$   
 $= 10011010101$

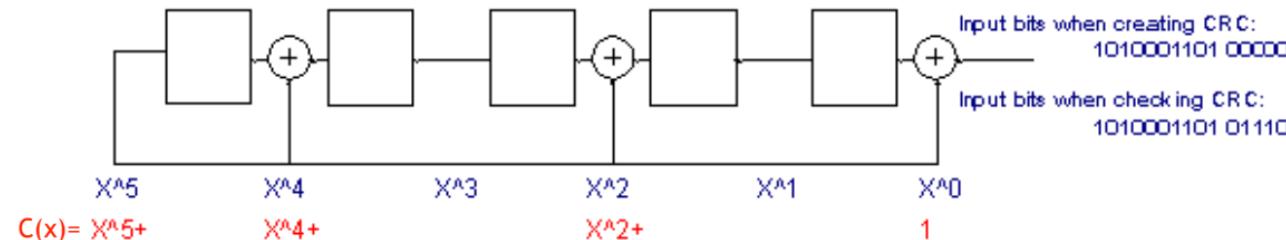


CRC Calculation using Polynomial Long Division

# Implementing the CRC algorithm in hardware

con poche XOR ed è molto economico i.e. circuito da costruire

- A  $r$ -long CRC can be easily computed by a circuit with  $r$  1-bit registers and some XOR ports (these circuits are embedded directly in network adapters)
- The circuit is implemented as follows:
  - The register contains  $r$  bits, equal to the length of the CRC
  - There are up to  $r$  XOR gates
  - The presence or absence of a gate corresponds to the presence or absence of a term in the divisor polynomial  $C(X)$
  - Bits are shifted in, from MSB to LSB
  - After all bits are shifted, the content of registers give the requested CRC
- Same circuit can be used to generate and to check the CRC
- Example: Message  $M=1010001101$ ;  $C=110101$ ;  $r=5$ ; resulting CRC:  $R=01110$



# Cyclic Redundancy Check (CRC)

## Properties of Generator Polynomial

- Let  $P(x)$  represent what the sender sent  
and  $P'(x) = P(x) + E(x)$  is the received string.  
messaggio  
messaggio alterato
- $E(x) = x^i$  means that the  $i$ -th bit has been flipped in the corresponding position in the message  $P(x)$ .
- We know that  $P(x)/C(x)$  leaves a remainder of 0, but if  $E(x)/C(x)$  leaves a remainder of 0, then either  $E(x) = 0$  or  $C(x)$  is factor of  $E(x)$ .  
POL. GENERATORE
- When  $C(x)$  is a factor of  $E(x)$  we have problem; errors go unnoticed.
- If there is a single bit error then  $E(x) = x^i$ , where  $i$  determines the bit in error. If  $C(x)$  contains two or more terms it will never divide  $E(x) = x^i$ , so all single bit errors will be detected.  
perché se  $E$  ha 1 sola cifra e  $C$  ne ha più di una allora non si può fare la divisione

# Cyclic Redundancy Check (CRC)

→ Una codifica polinomiale con  $n$  bit di controllo può rilevare  
Sequenze di errori di lunghezza  $\leq n$

- In general, it is possible to prove that the following **types of errors** can be detected by a  $C(x)$  with the stated properties
  - **All single-bit errors**, if the  $x_k$  and  $x_0$  terms have nonzero coefficients.
  - **All double-bit errors**, if  $C(x)$  has a factor with at least three terms.
  - **Any odd number of errors**, if  $C(x)$  contains the factor  $(x+1)$ .  
Se  $E(x)$  contiene un numero dispari di termini (errori) allora non potrà mai essere diviso per  $(x+1)$  nel sistema a modulo 2.
  - Any “burst” error (i.e., sequence of consecutive error bits) for which the length of the burst is less than  $k$  bits. (Most burst errors of larger than  $k$  bits can also be detected.)
- A polynomial  $C(x)$  can be factored as  $C(x)=C_1(x)*...*C_n(x)$ , so if a message can be divided by  $C(x)$ , it means that it can be divided by each  $C_1(x), \dots, C_n(x)$   
P<sup>per</sup> quindi
- hence, if we want to ensure several properties *at the same time*, it is enough to use a generator polynomial which is the product of the polynomials representing each property

# Cyclic Redundancy Check (CRC)

- Six generator polynomials that have become international standards are:
  - $\text{CRC-8} = x^8 + x^2 + x + 1$
  - $\text{CRC-10} = x^{10} + x^9 + x^5 + x^4 + x + 1$
  - $\text{CRC-12} = x^{12} + x^{11} + x^3 + x^2 + x + 1$
  - **CRC-16** =  $x^{16} + x^{15} + x^2 + 1$
  - $\text{CRC-CCITT} = x^{16} + x^{12} + x^5 + 1$
  - **CRC-32** =  $x^{32} + x^{26} + x^{23} + x^{22} + x^{16} + x^{12} + x^{11} + x^{10} + x^8 + x^7 + x^5 + x^4 + x^2 + x + 1$
- **CRC-32** is used in **802.3** (Ethernet), **802.11** (WiFi), ...
- PPP can use **CRC-32** or **CRC-16** with a different polynomial
- For many other examples see [https://en.wikipedia.org/wiki/Polynomial\\_representations\\_of\\_cyclic\\_redundancy\\_checks](https://en.wikipedia.org/wiki/Polynomial_representations_of_cyclic_redundancy_checks)

# Reliable Transmission

→ il più usato a liv. D.L. per rilevare errori.

- **CRC** is used to *detect* errors, not to correct them.
  - Some error codes are strong enough to correct errors up to a given limit (e.g. Hamming), but the overhead is typically too high.
- Corrupt frames must be discarded
  - Lo piuttosto che correggere è più economico richiedere l'invio, anche se non è sempre possibile farlo (es. satellite)
- An *unreliable* link-level protocol leaves to higher level protocols the decision to retransmit messages, if needed.
  - e.g. Ethernet, WiFi
- A *reliable* link-level protocol must recover from these discarded frames.
  - e.g. PPP
  - Obtained using a combination of two fundamental mechanisms:  
Acknowledgements and Timeouts

## Reliable Transmission

FRAME

messaggio che porta informazioni di controllo, "i dati che mi hai mandato fin'ora sono giusti"

- An **acknowledgement (ACK)** is a small control frame that a protocol sends back to its peer saying that it has received the earlier frame.
  - (A *control frame* is a frame with header only - no data).
- The receipt of an acknowledgement indicates to the sender of the original frame that its frame was successfully delivered.
- If the sender does not receive an acknowledgment after a reasonable amount of time, then it retransmits the original frame.
- The action of “waiting a reasonable amount of time” is called a **timeout**.
- The general strategy of using acknowledgements and timeouts to implement reliable delivery is sometimes called **Automatic Repeat reQuest (ARQ)**.

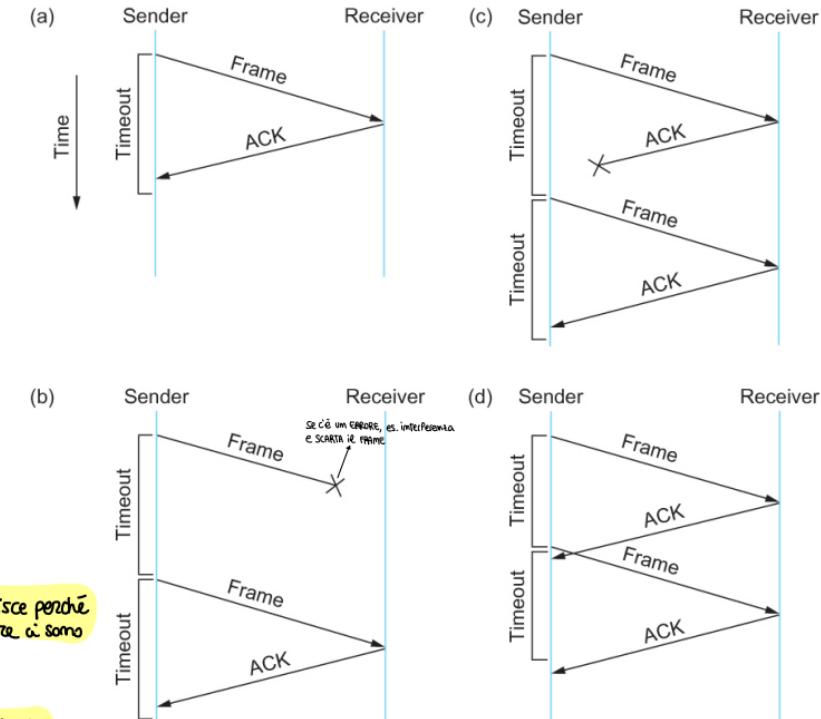
## Stop and Wait Protocol

- Idea of stop-and-wait protocol is straightforward
    - After transmitting one frame, the sender waits for an acknowledgement before transmitting the next frame.
    - If the acknowledgement does not arrive after a certain period of time, the sender times out and retransmits the original frame

- Timeline showing four different scenarios for the stop-and-wait algorithm.

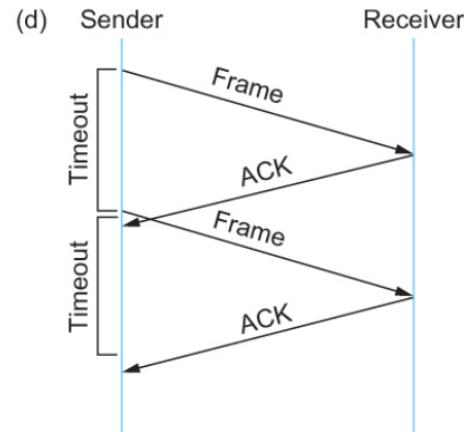
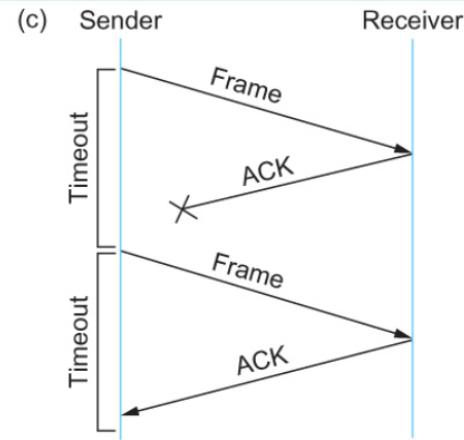
- (a) The ACK is received before the timer expires;
  - (b) the original frame is lost;
  - (c) the ACK is lost; → così riceve 2 copie dello stesso pacchetto, come cap effettivamente me ha ricevuti 2? erano volontari oppure
  - (d) the timeout fires too soon → Stai evitare di commissione?

- Cases (c), (d) are problematic



# Stop and Wait Protocol

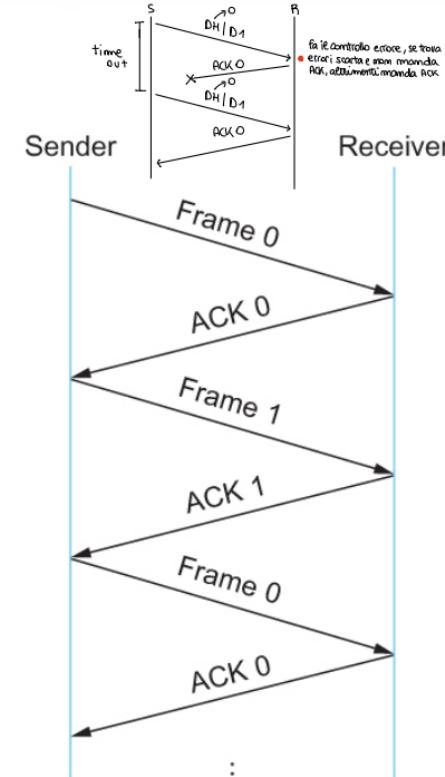
- If the acknowledgment is lost or delayed in arriving
    - The sender times out and retransmits the original frame, but the receiver will think that it is the next frame since it has correctly received and acknowledged the first frame
    - As a result, duplicate copies of frames will be delivered



# Stop and Wait Protocol → é um sistema perto

- How to solve this

- Use 1 bit sequence number (0 or 1)
  - Each frame is numbered (either 0 or 1)
    - This information is put in the header
  - In case of timeout, a frame is retransmitted with the same number
  - When the sender retransmits frame 0, the receiver can determine that it is seeing a second copy of frame 0 rather than the first copy of frame 1 and therefore can ignore it (the receiver still acknowledges it, in case the first acknowledgement was lost)



## Stop and Wait Protocol

- The sender has only one outstanding frame on the link at a time
    - This may be far below the link's capacity
  - Consider a 1.5 Mbps link with a 45 ms RTT
    - The link has a delay  $\times$  bandwidth product of  $67.5 \text{ Kb} \approx 8 \text{ KB}$
    - Since the sender can send only one frame per RTT and assuming a frame size of 1 KB, the maximum Sending rate is:
      - Bits per frame  $\div$  Time per frame  $= 1024 \times 8 \div 0.045 = 182 \text{ Kbps}$
      - About one-eighth of the link's capacity
      - The longer the RTT, the less we use the link
    - To use the link fully, then sender should transmit up to eight frames before having to wait for an acknowledgement

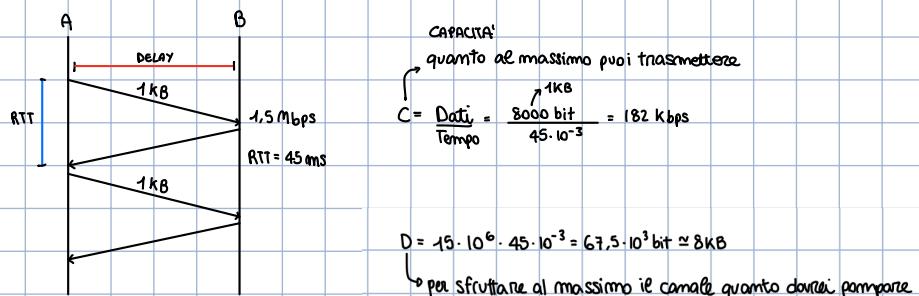
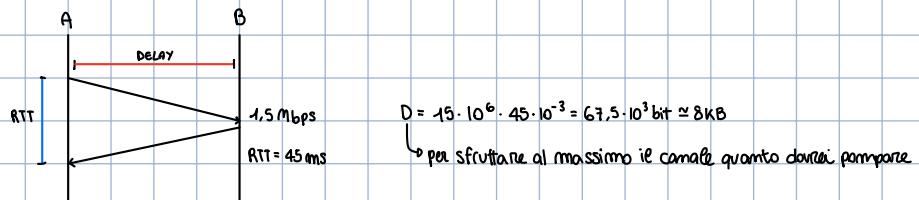
10 tempo ANDATA - RITORN

→ quantità max di dati che posso mandare prima che il primo pacchetto corrispondente arrivi

→  $D = 15 \cdot 10^6 \cdot 45 \cdot 10^{-3} = 67,5 \cdot 10^3 \text{ bit} \approx 8 \text{ KB}$   
 per sfruttare al massimo il canale quanto dati si possono

Se posso inviare solo un frame vuol dire che sto andando a 182 kbps, contro gli 1,5Mbps PREVISTI !! SPRECO

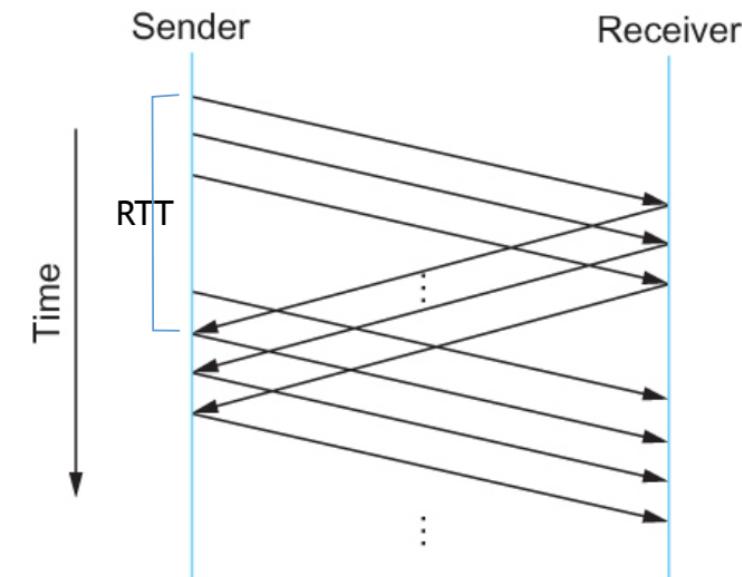
## STOP & WAIT (poco efficiente)



# Sliding Window Protocol

→ usato nel TCP

- The Sliding Window Protocol is a variant (better, a generalisation) of stop-and-wait protocol
- We allow the sender to send many frames during a *round* (RTT), even before receiving the first ACK, in order to fill the link capacity
- (Used sometimes at link level, more often at transport level)



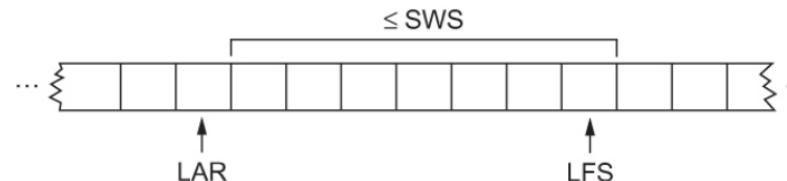
oltre ai bit abbiamo anche un numero di sequenza.

Timeline for Sliding Window Protocol

# Sliding Window Protocol

- Sender assigns a sequence number denoted as *SeqNum* to each frame.
  - Assume it can grow infinitely large
- *Sender* maintains three variables
  - *Sending Window Size (SWS)*: Upper bound on the number of outstanding (unacknowledged) frames that the sender can transmit
  - *Last Acknowledgement Received (LAR)*: Sequence number of the last acknowledgement received
    - ↳ # dell'ultimo frame inviato e ricevuto correttamente dal receiver (ovvero l'ultimo ACK ricevuto)
  - *Last Frame Sent (LFS)*: Sequence number of the last frame sent
    - ↳ # dell'ultimo frame che ho spedito
- Sender also maintains the following invariant

$$\text{LFS} - \text{LAR} \leq \text{SWS}$$

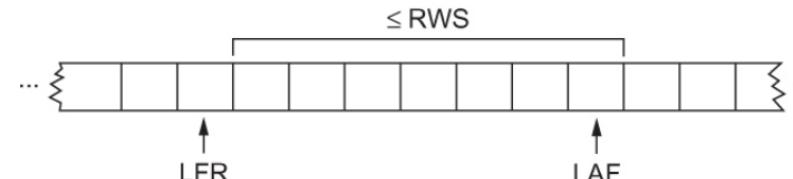


## Sliding Window Protocol

- When an acknowledgement arrives, the sender moves LAR to right, then goes in condition to transmit another frame
- Also, the sender associates a timer with each frame it transmits
- It retransmits the frame (with the same sequence number) if the timer expires before the ACK is received
- Note that the sender has to be willing to buffer up to SWS frames (why?)

# Sliding Window Protocol

- Receiver maintains three variables
- **Receiving Window Size (RWS)**: Upper bound on the number of out-of-order frames that the receiver is willing to accept
  - **Largest Acceptable Frame (LAF)**: Largest acceptable frame sequence number  
↳ upperbound del sequence number che può accettare
  - **Last Frame Received (LFR)**: Sequence number of the last frame received  
↳ fino a LFR i frame sono arrivati correttamente
- Receiver also maintains the following invariant  
**LAF - LFR  $\leq$  RWS**



# Sliding Window Protocol



- When a frame with sequence number SeqNum arrives, what does the receiver do?
  - If  $LFR < \text{SeqNum} \leq LAF \Rightarrow$  Accept it
  - If  $\text{SeqNum} \leq LFR$  or  $\text{SeqNum} > LAF \Rightarrow$  Discard it (the frame is outside the receiver window)
- Now the receiver needs to decide whether or not to send an ACK
  - Let  $\text{SeqNumToAck}$  denote the largest sequence number not yet acknowledged, such that all frames with sequence number less than or equal to  $\text{SeqNumToAck}$  have been received → avviso mandando un ack fino al primo frame che mi manca, anche se magari ho già quelli dopo
  - cumulative acknowledgement:** the receiver acknowledges the receipt until  $\text{SeqNumToAck}$  even if high-numbered packets have been received
  - In this way, we can buffer out-of-order frames, and acknowledge them cumulatively
- The receiver then sets  $LFR = \text{SeqNumToAck}$  and shifts window:  $LAF = LFR + RWS$

## Sliding Window Protocol

- For example, suppose LFR = 5 and RWS = 4  
(i.e. the last ACK that the receiver sent was for seq. no. 5)
  - LAF = 9
- If frames 7 and 8 arrive, they will be buffered because they are within the receiver window
  - But no ACK will be sent since frame 6 is yet to arrive
  - Frames 7 and 8 are out of order
- Frame 6 arrives (it is late because it was lost first time and had to be retransmitted)
- Now Receiver:
  - acknowledges Frame 8
  - bumps LFR to 8
  - bumps LAF to 12

## Issues with Sliding Window Protocol

- When timeout occurs, the amount of data in transit decreases
  - Since the sender is unable to advance its window
- When the packet loss occurs, this scheme is no longer keeping the pipe full
  - The longer it takes to notice that a packet loss has occurred, the more severe the problem becomes
- How to improve this
  - Negative Acknowledgement (NAK)
  - Additional Acknowledgement
  - Selective Acknowledgement

# Issues with Sliding Window Protocol

- **Negative Acknowledgement (NAK)**

- Receiver sends NAK for frame 6 when frame 7 arrives (in the previous example)
- However this is unnecessary since sender's timeout mechanism will be sufficient to catch the situation

- **Additional Acknowledgement**

- Receiver sends additional ACK for frame 5 when frame 7 arrives
- Sender uses duplicate ACK as a clue for frame loss
- Used in TCP

- **Selective Acknowledgement**

- Receiver will acknowledge exactly those frames it has received, rather than the highest number frames
- Receiver will acknowledge frames 7 and 8
- Sender knows frame 6 is lost
- Sender can keep the pipe full (at the price of additional complexity)

# Issues with Sliding Window Protocol

How to select the window size?

→ quantità di dati per riempire il tubo (il max che ci sta prima di un ACK)

- **SWS** is easy to compute:  $\text{Delay} \times \text{Bandwidth}$
- RWS can be anything → finestra dei frame non ancora "acknowledgiati"
- Two common setting
  - RWS = 1: No buffer at the receiver for frames that arrive out of order
  - RWS = SWS: The receiver can buffer frames that the sender transmits
- If RWS > SWS, the receiver buffer will be never filled
- If RWS < SWS, then sender could send frames which cannot be buffered and must be discarded and retransmitted later
- The case SWS=RWS=1 is the stop-and-wait ARQ.

# Issues with Sliding Window Protocol

- Finite Sequence Number
  - Frame sequence number is specified in the header field
  - Finite size
    - e.g., with 3 bit, eight possible sequence number: 0, 1, 2, 3, 4, 5, 6, 7
  - It is necessary to wrap around
  - How to distinguish between different incarnations of the same sequence number?
  - Number of possible sequence number must be larger than the number of outstanding frames allowed
    - Stop and Wait: One outstanding frame
      - 2 distinct sequence number (0 and 1)
    - Sliding window: Let  $\text{MaxSeqNum}$  be the maximum sequence number (usually  $=2^n-1$  where  $n$  is the number of bits in the header field)
      - $\text{SWS} \leq \text{MaxSeqNum}$
      - Is this sufficient?

# Issues with Sliding Window Protocol

- $SWS \leq \text{MaxSeqNum}$ 
  - Is this sufficient? Depends on RWS
    - If  $RWS = 1$ , then sufficient
    - If  $RWS = SWS$ , then not good enough
- For example, we have eight sequence numbers,  $\text{MaxSeqNum}=7$   
0, 1, 2, 3, 4, 5, 6, 7, and  $RWS = SWS = 7$ 
  - Sender sends frames 0, 1, ..., 6.
  - Receiver receives frames 0, 1, ..., 6 and sends ACKs 0, 1, ..., 6; now it is waiting for frames 7, 0, ..., 5.
  - But all the ACKs are lost! So the sender retransmits 0, 1, ..., 6.
  - Receiver accepts the frames 0, 1, ..., 5, which are *mistaken for the new ones*, and put in the buffer **after** the slot for 7!

# Issues with Sliding Window Protocol

- To avoid this: if  $RWS = SWS$  then it must be  
$$SWS \leq (\text{MaxSeqNum} + 1)/2 = 2^{n-1}$$
where  $n$  are the bits for the sequence number
- In this way, only half sequence numbers can be outstanding at any time
- In the example above, let us set  $RWS = SWS = 4$ 
  - Sender sends frames 0, 1, 2, 3, and stops waiting for acks
  - Receiver receives frames 0, 1, 2, 3 and sends ACKs 0, 1, 2, 3; now it is waiting for frames 4, 5, 6, 7 (it has up to 4 free slots).
  - But all the ACKs are lost, so the sender retransmits 0, 1, 2, 3.
  - Receiver was waiting for 4,...,7, so it recognises that the frames 0, 1, 2, 3 just received are copies of old ones: he does not accept these copies, but sends the acks again.

## Summary about Sliding Window Protocol

- Serves three different aims
  - Reliable → è AFFIDABILE, i messaggi arrivano
  - Preserve the order → nel caso dovesse arrivare in disordine, li riordina.
    - Each frame has a sequence number
    - The receiver makes sure that it does not pass a frame up to the next higher-level protocol until it has already passed up all frames with a smaller sequence number
  - Frame control
    - Receiver is able to throttle the sender
    - Keeps the sender from overrunning the receiver from transmitting more data than the receiver is able to process